



Maximally accurate

Maximally specific

espresso

2.23192

coffee

2.19914

beverage

1.93214

liquid

1.89367

fluid

1.85519

# Brewing Deep Networks With Caffe

Look for  
additional notes  
for the slides



Yangqing Jia

# What is, and Why Caffe?

- Pure C++/CUDA Implementation
- Fast, well-tested code
- Tools, demos, and recipes
- Seamless switch between CPU and GPU
  - `Caffe::set_mode(Caffe::GPU);`



Prototype



Training



Deployment

# Statistics...

- Speed with Krizhevsky's 2012 model:
  - K40 / Titan: **2 ms/image**, K20: 2.6ms
  - **(40 million images / day)**
  - 8-core CPU: ~20 ms/image
- **~ 8K** lines of C/C++ code
  - with unit test: ~14k

● C++ 84.4%

● Python 10.7%

● Cuda 3.5%

● Other 1.4%

\* Not counting image I/O time. Details at [http://caffe.berkeleyvision.org/performance\\_hardware.html](http://caffe.berkeleyvision.org/performance_hardware.html)

# Do I want Caffe If I...

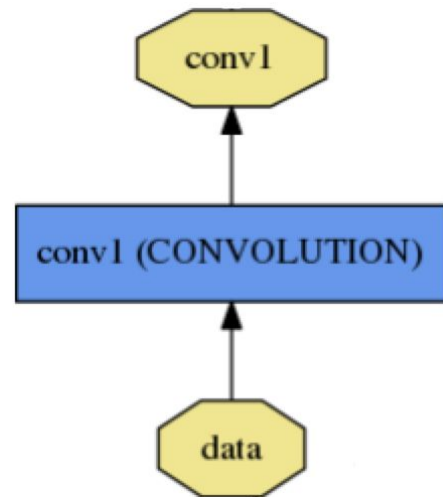
- Have small or medium scale applications?
  - Scripting languages may save engineering time indeed.
- Prefer simpler scripting languages?
  - We now provide Python and Matlab wrappers.
- Hate tricky compilation issues?
  - Recipes on Caffe webpage, and github.
  - Virtualbox / EC2 images to be provided soon.

# A Caffe Layer

```
name: "conv1"  
type: CONVOLUTION  
bottom: "data"  
top: "conv1"  
convolution_param {  
    num_output: 20  
    kernel_size: 5  
    stride: 1  
    weight_filler {  
        type: "xavier"  
    }  
}
```

name, type, and the  
connection structure  
(input blobs and  
output blobs)

layer-specific  
parameters



# A Caffe Network

- A network is a set of layers connected as a DAG:

**name:** "dummy-net"

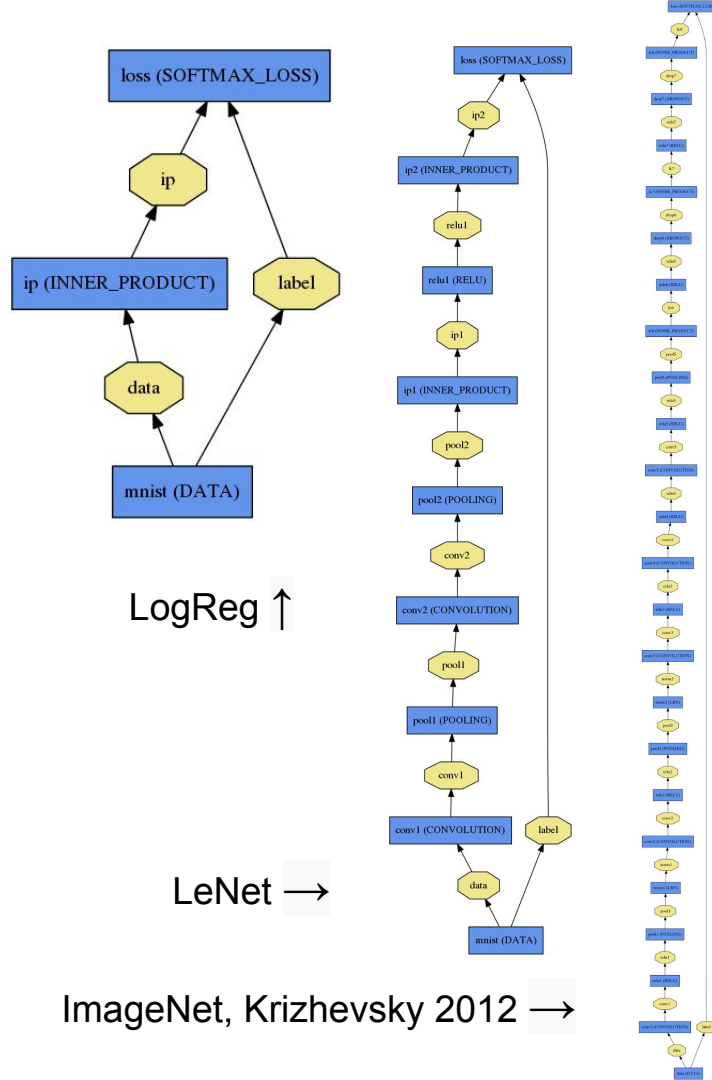
**layers** { name: "data" ...}

**layers** { name: "conv" ...}

**layers** { name: "pool" ...}

... more layers ...

**layers** { name: "loss" ...}



# Training a Caffe Net

Write a solver protobuffer:

```
train_net: "lenet_train.prototxt"  
base_lr: 0.01  
momentum: 0.9  
weight_decay: 0.0005  
max_iter: 10000  
snapshot_prefix: "lenet_snapshot"  
solver_mode: GPU
```

All you need to run things on the GPU.

# End to End Recipe...

- Convert the data to Caffe-format
  - leveldb, hdf5/.mat, list of images, LMDB, etc.
- Write a Network Definition
- Write a Solver Protobuffer text
- Train with the provided train\_net tool
  - `build/tools/train_net.bin solver.prototxt`
- Examples are your friends
  - `caffe/examples/mnist, cifar10, imagenet`
  - `caffe/tools/*.bin`

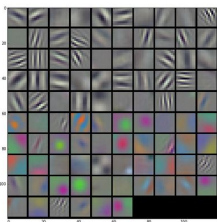


# Peeking into Networks



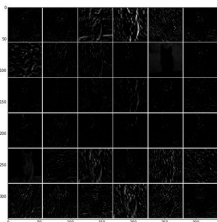
The first layer filters, conv1

```
In [8]: # the parameters are a list of [weights, biases]
filters = net.params['conv1'][0].data
vis_square(filters.transpose(0, 2, 3, 1))
```

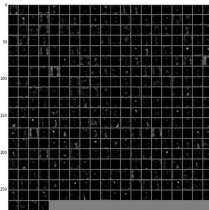


The first layer output, conv1 (rectified responses of the filters above, first 36 only)

```
In [9]: feat = net.blobs['conv1'].data[4, :36]
vis_square(feat, padval=1)
```

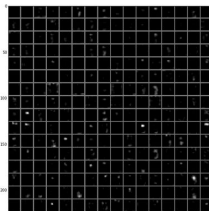


```
In [13]: feat = net.blobs['conv4'].data[4]
vis_square(feat, padval=0.5)
```



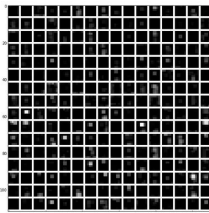
The fifth layer output, conv5 (rectified, all 256 channels)

```
In [14]: feat = net.blobs['conv5'].data[4]
vis_square(feat, padval=0.5)
```

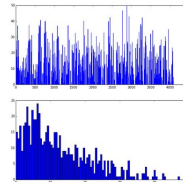


The fifth layer after pooling, pool5

```
In [15]: feat = net.blobs['pool5'].data[4]
vis_square(feat, padval=1)
```

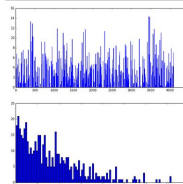


```
In [16]: feat = net.blobs['fc6'].data[4]
plt.subplot(2, 1, 1)
plt.plot(feat.flat)
plt.subplot(2, 1, 2)
_ = plt.hist(feat.flat[feat.flat > 0], bins=100)
```



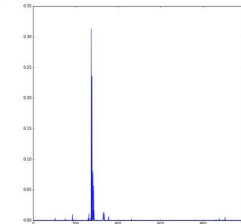
The second fully connected layer, fc7 (rectified)

```
In [17]: feat = net.blobs['fc7'].data[4]
plt.subplot(2, 1, 1)
plt.plot(feat.flat)
plt.subplot(2, 1, 2)
_ = plt.hist(feat.flat[feat.flat > 0], bins=100)
```



```
In [18]: feat = net.blobs['prob'].data[4]
plt.plot(feat.flat)
```

```
Out[18]: <matplotlib.lines.Line2D at 0x12b260710>
```



# A Quick Sip of Brewed Models

<http://demo.caffe.berkeleyvision.org/>

(demo code to be open-sourced soon)



Maximally accurate

Maximally specific

cat

1.80727

domestic cat

1.74727

feline

1.72787

tabby

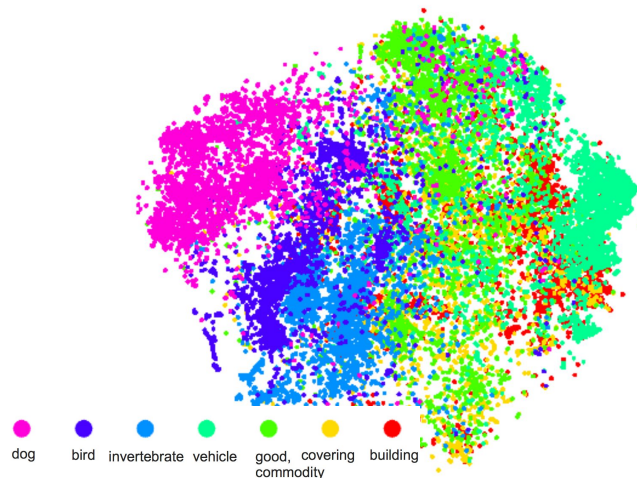
0.99133

domestic animal

0.78542

# Transfer Learned Knowledge

- Taking a pre-trained model and finetune it for related tasks [Zeiler-Fergus] [DeCAF] [OverFeat]



# Dogs vs Cats: top 10% in 10 minutes

- Simply change a few lines in the layer definition

<pre>layers {   name: "data"   type: DATA   data_param {     source: "ilsvrc12_train_leveldb"     mean_file: "../data/ilsvrc12/i     batch_size: 256     crop_size: 227     mirror: true   } }</pre>	→ ←	<pre>layers {   name: "data"   type: DATA   data_param {     source: "dogs-vs-cats-leveldb"     mean_file: "../data/ilsvrc1:     batch_size: 256     crop_size: 227     mirror: true   } }</pre>
<pre>layers {   name: "fc8"   type: INNER_PRODUCT   blobs_lr: 1   blobs_lr: 2   weight_decay: 1   weight_decay: 0   inner_product_param {     num_output: 1000   } }</pre>	→ ←	<pre>layers {   name: "fc8-dogcat"   type: INNER_PRODUCT   blobs_lr: 1   blobs_lr: 2   weight_decay: 1   weight_decay: 0   inner_product_param {     num_output: 2   } }</pre>

Input:  
A different source

Last Layer:  
A different classifier

# Dogs vs Cats: top 10% in 10 minutes

`build/tools/finetune_net.bin` `dogcat_solver.prototxt` `pretrained_imagenet_model`

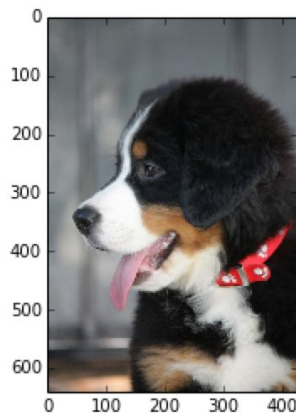
Under the hood (loosely speaking):

```
net = new Caffe::Net(  
    "dogcat_solver.prototxt");  
net.CopyTrainedNetFrom(  
    pretrained_model);  
solver.Solve(net);
```

Example code to be made available at  
[caffe/examples/dogs-vs-cats/](http://caffe/examples/dogs-vs-cats/)

```
plt.imshow(image)  
scores = net.predict([image]).flatten()  
if scores[1] > scores[0]:  
    print 'Woof it is a DOG!'  
else:  
    print 'Yiss it is a CAT!'
```

Woof it is a DOG!



# Object Detection

## R-CNN: Regions with Convolutional Neural Networks

<http://nbviewer.ipython.org/github/BVLC/caffe/blob/dev/examples/detection.ipynb>

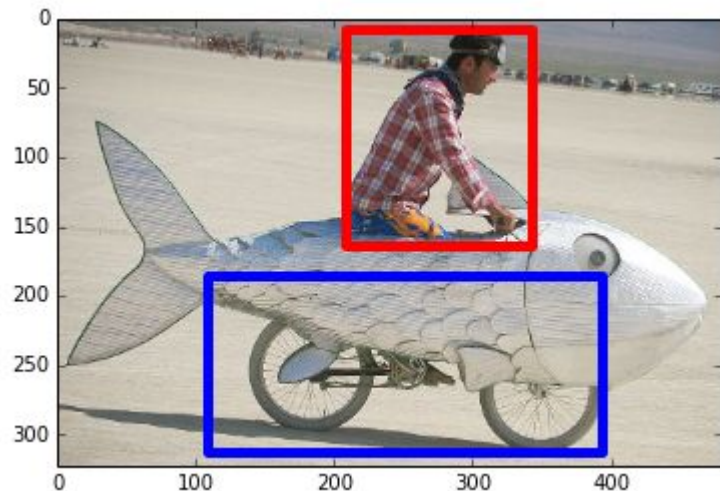
Full R-CNN scripts available at

<https://github.com/rbgirshick/rcnn>

Ross Girshick et al.

*Rich feature hierarchies for accurate object detection and semantic segmentation*

Oral Session 2A, Tue 1:30 pm



# Visual Style Recognition

Sergey Karayev, <http://vislab.berkeleyvision.org/>, demo available online

---

Ethereal



HDR



Melancholy



Minimal



Other Styles:

[Vintage](#)

[Long Exposure](#)

[Noir](#)

[Pastel](#)

[Macro](#)

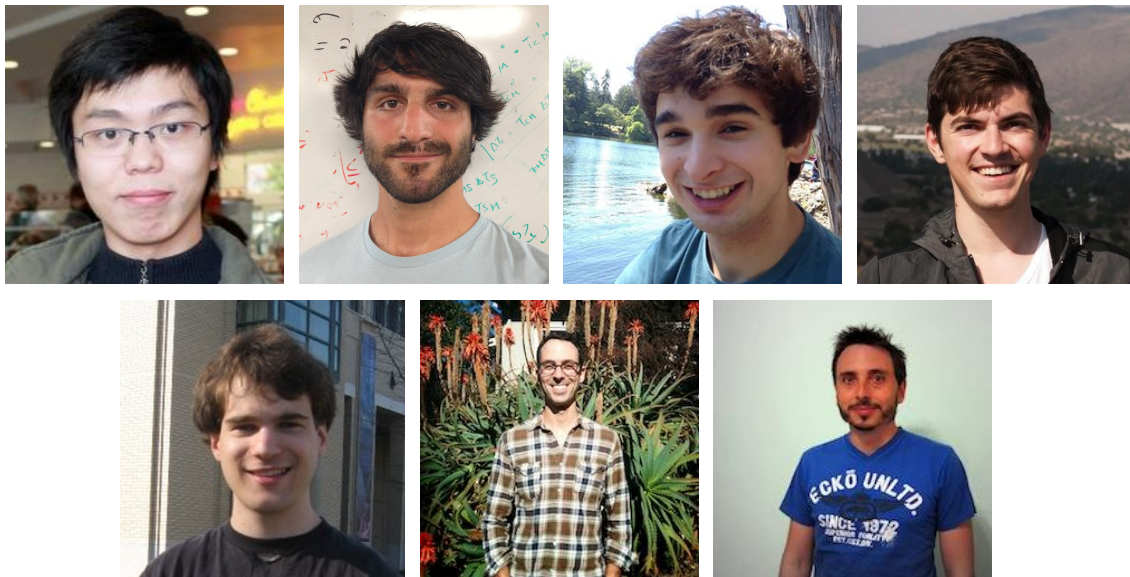
... and so on.

# In One Sip

Caffe...

- is C++/CUDA friendly
- is fast
- is state-of-the-art
- has tips, recipes, demos
- all available under an open-source initiative





Yangqing Jia, Evan Shelhamer, Jeff Donahue, Sergey Karayev  
Jonathan Long, Ross Girshick, Sergio Guadarrama

